

```
In [35]: import pandas as pd
import numpy as np
from pandas_profiling import ProfileReport
from scipy.stats import pearsonr
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.model_selection import GridSearchCV
# from xgboost import XGBClassifier
from sklearn.ensemble import AdaBoostClassifier
#from catboost import CatBoostRegressor
```

```
In [5]: !pip install xgboost
```

^C

```
In [16]: df = pd.read_csv('adult.csv')
```

```
In [137... df
```

```
Out[137...
   age  workclass  fnlwgt  education  education-  marital-  occupation  relationship  race
      0    39   State-gov   77516  Bachelors        13   Never-  Adm-  Not-in-family  White
      1    50  Self-emp-   83311  Bachelors        13   Married-  Exec-  Husband  White
      2    38   Private  215646   HS-grad         9   Divorced  Handlers-  Not-in-family  White
      3    53   Private  234721    11th         7   Married-  Handlers-  Husband  Black
      4    28   Private  338409  Bachelors        13   Married-  Prof-  Wife  Black
      ...    ...    ...    ...    ...    ...    ...    ...    ...    ...
  32556    27   Private  257302   Assoc-        12   Married-  Tech-  Wife  White
      32557    40   Private  154374   HS-grad         9   Married-  Machine-  Husband  White
      32558    58   Private  151910   HS-grad         9   Widowed  Adm-  Unmarried  White
      32559    22   Private  201490   HS-grad         9   Never-  Adm-  Own-child  White
      32560    52  Self-emp-  287927   HS-grad         9   Married-  Exec-  Wife  White
           inc
```

32561 rows × 15 columns

In [138]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   age                   32561 non-null  int64  
 1   workclass              32561 non-null  object  
 2   fnlwgt                 32561 non-null  int64  
 3   education              32561 non-null  object  
 4   education-num          32561 non-null  int64  
 5   marital-status         32561 non-null  object  
 6   occupation             32561 non-null  object  
 7   relationship           32561 non-null  object  
 8   race                   32561 non-null  object  
 9   sex                    32561 non-null  object  
10   capital-gain           32561 non-null  int64  
11   capital-loss           32561 non-null  int64  
12   hours-per-week         32561 non-null  int64  
13   country                 32561 non-null  object  
14   salary                 32561 non-null  object  
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

In [50]:

```
df.describe()
```

Out[50]:

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	40.437456
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	12.347429
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000

In [8]:

```
pf = ProfileReport(df)
```

In [9]:

```
pf.to_widgets()
```

In [15]:

```
display(df.corr().abs())
```

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week
--	-----	--------	---------------	--------------	--------------	----------------

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week
age	1.000000	0.076646	0.036527	0.077674	0.057775	0.068756
fnlwgt	0.076646	1.000000	0.043195	0.000432	0.010252	0.018768
education-num	0.036527	0.043195	1.000000	0.122630	0.079923	0.148123
capital-gain	0.077674	0.000432	0.122630	1.000000	0.031615	0.078409
capital-loss	0.057775	0.010252	0.079923	0.031615	1.000000	0.054256
hours-per-week	0.068756	0.018768	0.148123	0.078409	0.054256	1.000000

```
In [14]: (df.values.astype(str)=='?').sum()
```

Out[14]: 0

```
In [20]: n= df.shape[0]
n
```

Out[20]: 32561

```
In [17]: from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()
df['workclass'] = le.fit_transform(df['workclass'])
df['marital-status'] = le.fit_transform(df['marital-status'])
df['occupation'] = le.fit_transform(df['occupation'])
df['relationship'] = le.fit_transform(df['relationship'])
df['race'] = le.fit_transform(df['race'])
df['sex'] = le.fit_transform(df['sex'])
df['country'] = le.fit_transform(df['country'])
df['salary'] = le.fit_transform(df['salary'])
df['education'] = le.fit_transform(df['education'])

df.head()
```

Out[17]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex
0	39	7	77516	9	13	4	1	1	4	...
1	50	6	83311	9	13	2	4	0	4	...
2	38	4	215646	11	9	0	6	1	4	...
3	53	4	234721	1	7	2	6	0	2	...
4	28	4	338409	9	13	2	10	5	2	...

```
In [18]: x = df.drop(['salary'], axis = 1)
y = df['salary']
```

```
In [19]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size = 0.2, ...)
```

```
In [147... from sklearn.naive_bayes import GaussianNB
gb = GaussianNB()
gb.fit(x_train,y_train)
```

Out[147... GaussianNB()

```
In [36]: from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
```

```
In [149... y_pred = gb.predict(x_test)
print(classification_report(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))
print(accuracy_score(y_test,y_pred)*100)
```

	precision	recall	f1-score	support
0	0.81	0.94	0.87	4920
1	0.64	0.32	0.43	1593
accuracy			0.79	6513
macro avg	0.73	0.63	0.65	6513
weighted avg	0.77	0.79	0.76	6513

```
[[4635 285]
 [1082 511]]
79.01120835252571
```

```
In [121... rf = RandomForestClassifier()
rf.fit(x_train,y_train)
```

Out[121... RandomForestClassifier()

```
In [122... y_pred = rf.predict(x_test)
print(classification_report(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))
print(accuracy_score(y_test,y_pred)*100)
```

	precision	recall	f1-score	support
0	0.89	0.93	0.91	4920
1	0.74	0.63	0.68	1593
accuracy			0.86	6513
macro avg	0.81	0.78	0.79	6513
weighted avg	0.85	0.86	0.85	6513

```
[[4564 356]
 [ 587 1006]]
85.52126516198373
```

```
In [186... param_grid = {
    'n_estimators': [350,400],
    'max_features': ['auto', 'sqrt'],
    'max_depth' : [7,8,9,10],
    'criterion' :['gini', 'entropy']
}
```

```
In [ ]: CV_rfc = GridSearchCV(estimator=rf, param_grid=param_grid, cv= 3)
CV_rfc.fit(x_train, y_train)
```

```
In [155... CV_rfc.best_params_
```

```
Out[155... {'criterion': 'gini',
           'max_depth': 7,
           'max_features': 'sqrt',
           'n_estimators': 120}
```

```
In [221... rfcl=RandomForestClassifier(random_state=42, max_features='sqrt', n_estima
```

```
In [222... rfcl.fit(x_train, y_train)
```

```
Out[222... RandomForestClassifier(max_depth=22, max_features='sqrt', n_estimators=320,
                               random_state=42)
```

```
In [223... pred=rfcl.predict(x_test)
```

```
In [225... print("Accuracy for Random Forest on CV data: ",accuracy_score(y_test,pred)

Accuracy for Random Forest on CV data:  86.3196683555965
```

```
In [20]: ab_clf = AdaBoostClassifier(random_state=42)
```

```
In [21]: parameters = {
           'n_estimators': [2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 20],
           'learning_rate': [(0.97 + x / 100) for x in range(0, 8)],
           'algorithm': ['SAMME', 'SAMME.R']
       }
```

```
In [24]: clf = GridSearchCV(ab_clf, parameters, cv=5,n_jobs=-1)
clf.fit(x_train, y_train)
```

```
Out[24]: GridSearchCV(cv=5, estimator=AdaBoostClassifier(random_state=42), n_jobs=-
1,
           param_grid={'algorithm': ['SAMME', 'SAMME.R'],
                       'learning_rate': [0.97, 0.98, 0.99, 1.0, 1.01, 1.0
2,
                               1.03, 1.04],
                       'n_estimators': [2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 1
2,
                               20]})
```

```
In [26]: clf.best_params_
```

```
Out[26]: {'algorithm': 'SAMME.R', 'learning_rate': 0.98, 'n_estimators': 20}
```

```
In [50]: ab_clf1 = AdaBoostClassifier(random_state=42,algorithm= 'SAMME.R', learning
```

```
In [51]: ab_clf1.fit(x_train, y_train)
```

```
Out[51]: AdaBoostClassifier(learning_rate=0.98, n_estimators=320, random_state=42)
```

```
In [48]: predd=ab_clf1.predict(x_test)
```

```
In [49]: print("Accuracy for AdaBoostClassifier on CV data: ",accuracy_score(y_test
```

```
Accuracy for AdaBoostClassifier on CV data:  86.78028558268079
```

```
In [ ]:
```